

```
In [1]: import pandas as pd

In [2]: data = pd.read_csv('bank_note_data.csv')

Check the head of the Data

In [3]: data.head()

Out[3]:
   Image.Var  Image.Skew  Image.Curt  Entropy  Class
0    3.62160    8.6661  -2.8073 -0.44699    0
1    4.54590    8.1674  -2.4586 -1.46210    0
2    3.86600   -2.6383   1.9242  0.10645    0
3    3.45660    9.5228  -4.0112 -3.59440    0
4    0.32924   -4.4552   4.5718 -0.98880    0

In [4]: import seaborn as sns
%matplotlib inline

In [5]: sns.countplot(x='Class',data=data)

Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x26bb34edda0>

In [6]: sns.pairplot(data,hue='Class')

C:\Users\Marcial\Anaconda3\lib\site-packages\statsmodels\statsmodels\kde.py:494: RuntimeWarning: invalid value encountered in true_divide
  binned = fast_linbin(X,a,b,gridsize)/(delta*nobs)
C:\Users\Marcial\Anaconda3\lib\site-packages\statsmodels\statsmodels\kdetools.py:34: RuntimeWarning: invalid value encountered in double_scalars
  FAC1 = 2*(np.pi**bw/RANGE)**2
C:\Users\Marcial\Anaconda3\lib\site-packages\numpy\core\_methods.py:26: RuntimeWarning: invalid value encountered in reduce
  return umr_maximum(a, axis, None, out, keepdims)

Out[6]: <seaborn.axisgrid.PairGrid at 0x26bb34c9da0>

In [7]: from sklearn.preprocessing import StandardScaler

In [8]: scaler = StandardScaler()

In [9]: scaler.fit(data.drop('Class',axis=1))

Out[9]: StandardScaler(copy=True, with_mean=True, with_std=True)

In [10]: scaled_features = scaler.fit_transform(data.drop('Class',axis=1))

In [11]: df_feat = pd.DataFrame(scaled_features,columns=data.columns[:-1])
df_feat.head()

Out[11]:
   Image.Var  Image.Skew  Image.Curt  Entropy
0    1.121806    1.149455 -0.975970  0.354561
1    1.447066    1.064453 -0.895036 -0.128767
2    1.207810   -0.777352  0.122218  0.618073
3    1.063742    1.295478 -1.255397 -1.144029
4   -0.036772   -1.087038  0.736730  0.096587

In [12]: X = df_feat

In [13]: y = data['Class']

In [14]: from sklearn.model_selection import train_test_split

In [15]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

In [16]: import tensorflow as tf

C:\Users\Marcial\Anaconda3\lib\site-packages\h5py\_init_.py:34: FutureWarning: Conversion of the second argument of 'issubdtype' from 'float' to 'np.floating' is deprecated. In future, it will be treated as np.float64 == np.dtype(float).type.
  from .conv import register_converters as _register_converters

In [17]: df_feat.columns

Out[17]: Index(['Image.Var', 'Image.Skew', 'Image.Curt', 'Entropy'], dtype='object')

In [18]: image_var = tf.feature_column.numeric_column("Image.Var")
image_skew = tf.feature_column.numeric_column('Image.Skew')
image_curt = tf.feature_column.numeric_column('Image.Curt')
entropy =tf.feature_column.numeric_column('Entropy')

In [19]: feat_cols = [image_var,image_skew,image_curt,entropy]

In [20]: classifier = tf.estimator.DNNClassifier(hidden_units=[10, 20, 10], n_classes=2,feature_columns=feat_cols)

INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory: C:\Users\Marcial\AppData\Local\Temp\tmpw8v7z_z6
INFO:tensorflow:Using config: {'_model_dir': 'C:\Users\Marcial\AppData\Local\Temp\tmpw8v7z_z6', '_tf_random_seed': None, '_save_summary_steps': 100, '_save_checkpoints_step': None, '_save_checkpoints_secs': 600, '_session_config': None, '_keep_checkpoint_max': 5, '_keep_checkpoint_every_n_hours': 10000, '_log_step_count_steps': 100, '_train_distributed': None, '_device_fn': None, '_service': None, '_cluster_spec': <tensorflow.python.training.server_lib.ClusterSpec object at 0x000026BBA0F9FD0>, '_task_type': 'worker', '_task_id': 0, '_global_id_in_cluster': 0, '_master': '', '_evaluation_master': '', '_is_chief': True, '_num_ps_replicas': 0, '_num_worker_replicas': 1}

In [21]: input_func = tf.estimator.inputs.pandas_input_fn(x=X_train,y=y_train,batch_size=20,shuffle=True)

In [22]: classifier.train(input_fn=input_func,steps=500)

INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Saving checkpoints for 0 into C:\Users\Marcial\AppData\Local\Temp\tmpw8v7z_z6\model.ckpt.
INFO:tensorflow:loss = 13.792015, step = 1
INFO:tensorflow:Saving checkpoints for 48 into C:\Users\Marcial\AppData\Local\Temp\tmpw8v7z_z6\model.ckpt.
INFO:tensorflow:Loss for final step: 0.47980386.

Out[22]: <tensorflow.python.estimator.canned.dnn.DNNClassifier at 0x26bbacc7c18>

Model Evaluation

In [23]: pred_fn = tf.estimator.inputs.pandas_input_fn(x=X_test,batch_size=len(X_test),shuffle=False)

In [24]: note_predictions = list(classifier.predict(input_fn=pred_fn))

INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from C:\Users\Marcial\AppData\Local\Temp\tmpw8v7z_z6\model.ckpt-48
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.

In [25]: note_predictions[0]

Out[25]: {'class_ids': array([0], dtype=int64),
'classes': array([b'0'], dtype=object),
'logits': array([0.00157453], dtype=float32),
'logits': array([ 6.4522204], dtype=float32),
'probabilities': array([0.9984255, 0.00157453], dtype=float32)}

In [26]: final_preds = []
for pred in note_predictions:
    final_preds.append(pred['class_ids'][0])

In [27]: from sklearn.metrics import classification_report,confusion_matrix

In [28]: print(confusion_matrix(y_test,final_preds))

[[213  2]
 [ 10 187]]

In [29]: print(classification_report(y_test,final_preds))

              precision    recall  f1-score   support

    0             0.96         0.99         0.97         215
    1             0.99         0.95         0.97         197

avg / total          0.97         0.97         0.97         412

In [30]: from sklearn.ensemble import RandomForestClassifier

In [31]: rfc = RandomForestClassifier(n_estimators=200)

In [32]: rfc.fit(X_train,y_train)

Out[32]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_split=1e-07, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=200, n_jobs=1, oob_score=False, random_state=None,
verbose=0, warm_start=False)

In [33]: rfc_preds = rfc.predict(X_test)

In [34]: print(classification_report(y_test,rfc_preds))

              precision    recall  f1-score   support

    0             0.99         1.00         0.99         215
    1             0.99         0.99         0.99         197

avg / total          0.99         0.99         0.99         412

In [35]: print(confusion_matrix(y_test,rfc_preds))

[[214  1]
 [  2 195]]

It should have also done very well, possibly perfect! Hopefully you have seen the power of DNN!
```